

CCRE Executive Talking Points (5■Minute Script)

0:00–0:30 — Opening & Context

- We've modernized Treatment Selector by replacing legacy, scattered rule execution with **CCRE**—our unified, app■layer rule engine.
- Goal: one consistent decision model across channels, lower latency, and easier onboarding of new experiences.

0:30–1:15 — What Was Broken (Legacy)

- Legacy path relied on Adobe Target/legacy SQL rules and graph lookups owned in pockets.
- Result: duplicated logic, inconsistent behavior by channel, and higher end■to■end latency.
- Example: mobile app home latency dropped from **350+ ms** to **~250 ms** after we began shifting—**30–50%** reductions across contexts.

1:15–2:15 — What We Built (CCRE)

- CCRE runs rules in the app layer—not hidden in point tools.
- Single model for all rule execution; uniform consumption of APIs and databases via a shared data layer.
- Architecture blocks: **Planner & Optimizer** choose the cheapest evaluation path; **Caching & Logging** for speed/observability; the **Engine** executes; **TS** and **Context Enrichment** wrap channel needs.
- Channels now include Retail/WI, Email, Associate tools, IVR, and Self■Service—same rule once, everywhere.

2:15–3:15 — Why It's Better (Features & Numbers)

- Smart short■circuiting stops work early when outcomes are decided.
- Open, scalable architecture; **trace logs** for auditability.
- In■built **cost■aware optimization** selects the lowest■cost data source and operator order.
- Scale & latency: tested at **~20K RPS/region** with **p95 ≈ 35 ms**.
- Net effect: one consolidated rule plane that is faster, cheaper, and easier to operate.

3:15–4:00 — Business Outcomes (PZN & Onboarding)

- PZN expanded to all channels that need real■time data—no per■channel rewrites.
- Faster onboarding: Alts, Greenline, Mobile App Home (including the new youth app), plus additional WI use cases.
- Positioned to support WI's AEM migration (2026) without re■implementing rules per surface.

4:00–4:40 — Risk Reduction & Governance

- Single source of truth for decisions reduces drift and production variance.
- Full traceability: every decision has logs and inputs; simpler compliance reviews.
- Operational knobs—planner metrics, cache hit rates, and p95—all visible and tunable.

4:40–5:00 — The Ask / Next Steps

- Scale the unified model to remaining rule pockets and retire duplicative logic.
- Invest in rule authoring UX and self-service onboarding to accelerate adoption.
- Greenlight shared data-layer integrations with a few additional systems to unlock the next 20–30% latency/TCO gains.

Backup Q&A; One-Liners

- **Cost savings?** Fewer duplicate rule paths, less vendor overhead, and cost-based plan selection reduces expensive data calls.
- **Resilience?** Short-circuiting + caching reduce dependency pressure; a single hardened engine is easier to scale.
- **Governance?** Centralized rule logs, versioning, and channel-agnostic testing prevent regressions.
- **New channels?** Plug into Context Enrichment + TS; the rule stays the same.